

Fundamentals of Data Engineering

Week 07 - sync session

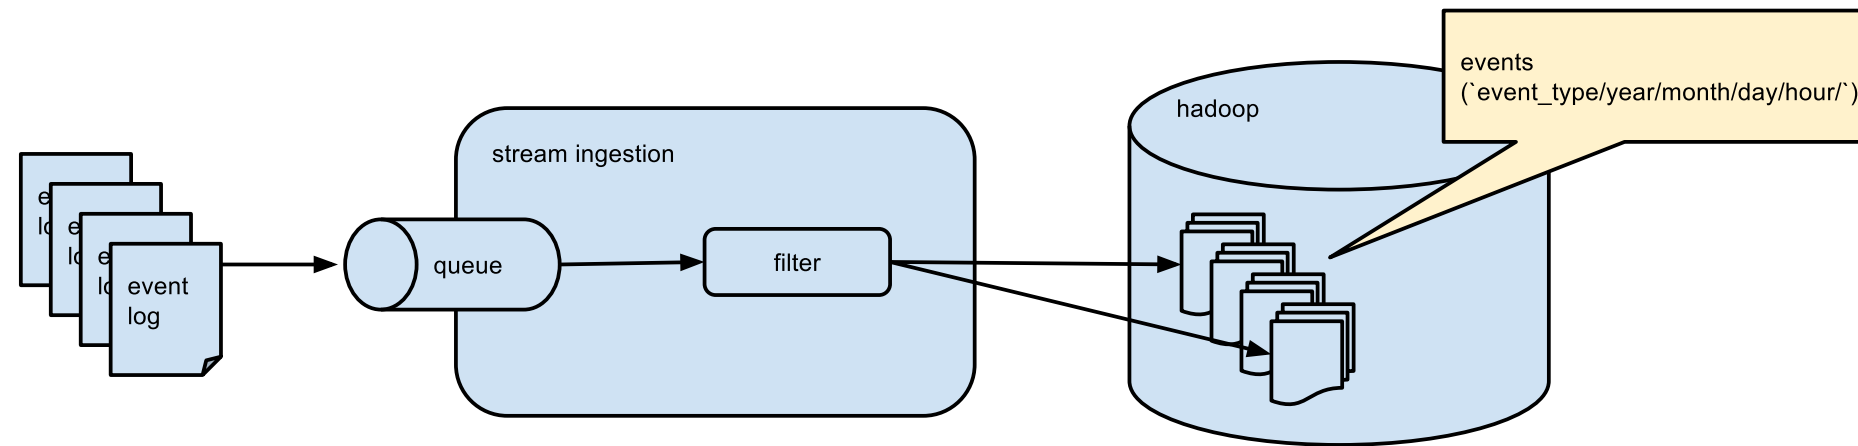
datascience@berkeley

While we're getting started

- Review your Assignment 06
- Get ready to share

::: notes Breakout at about 5 after the hour: - Check in with each group - have students share screen :::

Due Friday (PR)



::: notes - Last week we published and consumed messages with kafka - Now we'll consume messages

Spark Stack with Kafka

Setup

docker-compose.yml

```
---
version: '2'
services:
  zookeeper:
    image: confluentinc/cp-zookeeper:latest
    environment:
      ZOOKEEPER_CLIENT_PORT: 32181
      ZOOKEEPER_TICK_TIME: 2000
    expose:
      - "2181"
      - "2888"
      - "32181"
      - "3888"
    extra_hosts:
      - "moby:127.0.0.1"
```

::: notes and save a data file.

:::

up

```
docker-compose up -d
```

```
docker-compose logs -f kafka
```

::: notes Now spin up the cluster

```
docker-compose up -d
```

and watch it come up

```
docker-compose logs -f kafka
```

when this looks like it's done, you can safely detach with
Ctrl-C.

:::

use it

create a topic

```
docker-compose exec kafka \  
  kafka-topics \  
    --create \  
    --topic foo \  
    --partitions 1 \  
    --replication-factor 1 \  
    --if-not-exists \  
    --zookeeper zookeeper:32181
```

::: notes First, create a topic foo

```
docker-compose exec kafka kafka-topics --create --topic  
foo --partitions 1 --replication-factor 1 --if-not-exists --  
zookeeper zookeeper:32181 :::
```

Should show

```
Created topic "foo".
```

Check the topic

```
docker-compose exec kafka \  
  kafka-topics \  
    --describe \  
    --topic foo \  
    --zookeeper zookeeper:32181
```

Should show

```
Topic:foo PartitionCount:1 ReplicationFactor:1 Configs:  
Topic: foo Partition: 0 Leader: 1 Replicas: 1 Isr: 1
```

Publish some stuff to kafka

```
docker-compose exec kafka \  
  bash -c "seq 42 | kafka-console-producer \  
    --request-required-acks 1 \  
    --broker-list kafka:29092 \  
    --topic foo && echo 'Produced 42 messages.'"
```

::: notes Use the kafka console producer to publish some test messages to that topic

```
docker-compose exec kafka bash -c "seq 42 | kafka-  
console-producer --request-required-acks 1 --broker-list  
kafka:29092 --topic foo && echo 'Produced 42  
messages.'"
```

:::

Should show

```
Produced 42 messages.
```

Run spark using the myspark container

```
docker-compose exec myspark pyspark
```

::: notes Spin up a pyspark process using the myspark container

```
docker-compose exec myspark pyspark
```

We have to add some kafka library dependencies on the cli for now. :::

read stuff from kafka

At the pyspark prompt,

```
numbers = spark \  
  .read \  
  .format("kafka") \  
  .option("kafka.bootstrap.servers", "kafka:29092") \  
  .option("subscribe","foo") \  
  .option("startingOffsets", "earliest") \  
  .option("endingOffsets", "latest") \  
  .load()
```

::: notes At the pyspark prompt,

read from kafka

numbers = spark

.read

.format("kafka")

.option("kafka.bootstrap.servers", "kafka:29092")

See the schema

```
numbers.printSchema()
```

Cast it as strings

```
numbers_as_strings=numbers.selectExpr("CAST(key AS STRING)", "CAST(va
```

::: notes cast it as strings (you can totally use INTs if you'd like)

```
numbers_as_strings=numbers.selectExpr("CAST(key  
AS STRING)", "CAST(value AS STRING)") :::
```

Take a look

```
numbers_as_strings.show()
```

```
numbers_as_strings.printSchema()
```

```
numbers_as_strings.count()
```

::: notes then you can exit pyspark using either `ctrl-d` or `exit()`.

```
numbers_as_strings.show()
```

```
numbers_as_strings.printSchema()
```

```
:::
```

down

```
docker-compose down
```

Spark stack with Kafka with "real" messages

::: notes :::

docker-compose.yml file

```
---
version: '2'
services:
  zookeeper:
    image: confluentinc/cp-zookeeper:latest
    environment:
      ZOOKEEPER_CLIENT_PORT: 32181
      ZOOKEEPER_TICK_TIME: 2000
    expose:
      - "2181"
      - "2888"
      - "32181"
      - "3888"
    #ports:
      #- "32181:32181"
    extra_hosts:
```

::: notes - See ~/w205/spark-with-kafka-json folder's
docker-compose.yml - may use this cluster, may have it
working w/o myhdfs container :::

Pull data

```
curl -L -o github-example-large.json https://goo.gl/Hr6erG
```


Spin up the cluster

```
docker-compose up -d
```

Watch it come up

```
docker-compose logs -f kafka
```

- Detach with Ctrl-C

::: notes when this looks like it's done, detach :::

use it

create a topic

```
docker-compose exec kafka \  
  kafka-topics \  
    --create \  
    --topic foo \  
    --partitions 1 \  
    --replication-factor 1 \  
    --if-not-exists \  
    --zookeeper zookeeper:32181
```

::: notes docker-compose exec kafka kafka-topics --
create --topic foo --partitions 1 --replication-factor 1 --if-
not-exists --zookeeper zookeeper:32181 :::

Should see something like

```
Created topic "foo".
```

Check the topic

```
docker-compose exec kafka \  
  kafka-topics \  
    --describe \  
    --topic foo \  
    --zookeeper zookeeper:32181
```

::: notes docker-compose exec kafka kafka-topics --
describe --topic foo --zookeeper zookeeper:32181 ::: ##

Should see something like

```
Topic:foo    PartitionCount:1    ReplicationFactor:1 Configs:  
Topic: foo   Partition: 0      Leader: 1      Replicas: 1   Isr: 1
```

Publish some stuff to kafka

Check out our messages

```
docker-compose exec mids bash -c "cat /w205/github-example-large.json"
docker-compose exec mids bash -c "cat /w205/github-example-large.json"
```

::: notes ugly 1st pretty print :::

Individual messages

```
docker-compose exec mids bash -c "cat /w205/github-example-large.json"
```

::: notes Go over | jq stuff :::

Publish some test messages to that topic with the kafka console producer

```
docker-compose exec mid5 \
  bash -c "cat /w205/github-example-large.json \
    | jq '.[0]' -c \
    | kafkacat -P -b kafka:29092 -t foo && echo 'Produced 100 messages'"
```

::: notes docker-compose exec mid5 bash -c "cat /w205/github-example-large.json | jq '.[0]' -c | kafkacat -P -b kafka:29092 -t foo && echo 'Produced 100 messages.'"

:::

Should see something like

```
Produced 100 messages.
```

Run spark using the myspark container

```
docker-compose exec myspark pyspark
```

::: notes Spin up a pyspark process using the myspark container

```
docker-compose exec myspark pyspark
```

We have to add some kafka library dependencies on the cli for now. :::

read stuff from kafka

At the pyspark prompt,

```
messages = spark \  
  .read \  
  .format("kafka") \  
  .option("kafka.bootstrap.servers", "kafka:29092") \  
  .option("subscribe","foo") \  
  .option("startingOffsets", "earliest") \  
  .option("endingOffsets", "latest") \  
  .load()
```

::: notes At the pyspark prompt,

read from kafka

messages = spark

.read

.format("kafka")

.option("kafka.bootstrap.servers", "kafka:29092")

See the schema

```
messages.printSchema()
```

Cast as strings

```
messages_as_strings=messages.selectExpr("CAST(key AS STRING)", "CAST"
```

::: notes cast it as strings (you can totally use INTs if you'd like)

```
messages_as_strings=messages.selectExpr("CAST(key  
AS STRING)", "CAST(value AS STRING)") :::
```

Take a look

```
messages_as_strings.show()
```

```
messages_as_strings.printSchema()
```

```
messages_as_strings.count()
```

::: notes then you can exit pyspark using either `ctrl-d` or `exit()`.

```
messages_as_strings.show()
```

```
messages_as_strings.printSchema()
```

```
messages_as_strings.count() :::
```


Unrolling json

```
messages_as_strings.select('value').take(1)
```

```
messages_as_strings.select('value').take(1)[0].value
```

```
import json
```

```
first_message=json.loads(messages_as_strings.select('value').take(1)
```

```
first_message
```

```
print(first_message['commit']['committer']['name'])
```

```
::: notes messages_as_strings.select('value').take(1)
```

```
messages_as_strings.select('value').take(1)[0].value >>>
```

```
import json >>>
```

```
first_message=json.loads(messages_as_strings.select('value').take(1)[0].value) >>> first_message >>>
```

```
print(first_message['commit']['committer']['name']) Nico
```

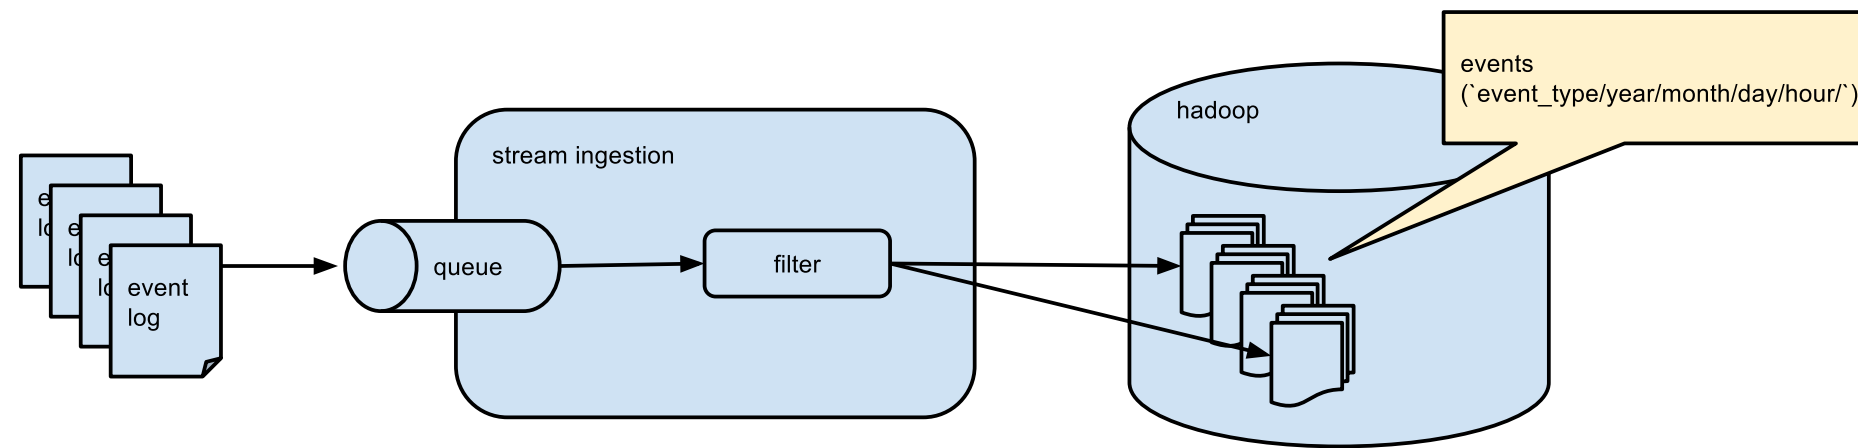
Down

```
docker-compose down
```

Assignment 07

- Step through this process using the Project 2 data
- What you turn in:
- In your /assignment-07-<user-name> repo:
- your docker-compose.yml
- once you've run the example on your terminal
 - Run `history > <user-name>-history.txt`
 - Save the relevant portion of your history as `<user-name>-annotations.md`
 - Annotate the file with explanations of what you were doing at each point (See `htmartin-annotations.md`)

Summary



::: notes - Review what we just did :::

Berkeley

SCHOOL OF
INFORMATION